

Saripalli Shanmukha Kiran Sagar

+91 6301287382 • shanmukhkiransagar@gmail.com • [linkedin.com/in/shanmukhkiransagar](https://www.linkedin.com/in/shanmukhkiransagar) • github.com/theoxfaber
AI / Software Engineering Intern | Rust • AI Systems • Backend | Open to 6–12-mo from Oct 2026 (Bangalore / Remote)

EDUCATION

Parul University — B.Tech, Artificial Intelligence & Data Science | Expected Graduation: 2028

Relevant coursework: Data Structures & Algorithms, DBMS, Operating Systems, Machine Learning, Probability & Statistics

EXPERIENCE

AI Coding Agent Evaluator — Alignerr (Labelbox) | Contract, Remote | Dec 2025 – May 2026

- Evaluated AI coding agents on realistic Git-repo tasks (debugging, refactoring, feature work, code review, testing) for implementation correctness, scope, and code quality.
- Assessed safety, honesty, instruction-following, verification, and engineering judgment; flagged edge cases, regressions, weak tests, and convention violations in generated diffs.

PROJECTS

Agentic Payment Risk Governor — Rust / Axum / React / MCP | [GitHub](#) • [Live Demo](#) • Razorpay AI Buildathon 2026

- Built a zero-trust proxy for agent-driven payments that prevents agents from accessing Razorpay keys and adjudicates every payment intent as ALLOW / REVIEW / BLOCK over authenticated REST/MCP.
- Solved lying-agent and collusion-ring problem with a four-plane pipeline: deterministic policy gates + anomaly detection (z-score, PSI drift) + sybil-graph union-find + calibrated logistic regression with Conformal Risk Control.
- Enforced fintech invariants: integer-paise math, fail-closed captured-balance checks, composite idempotency, pending-claim locks, receipt probing, constant-time auth, HMAC webhooks, tamper-evident SHA-256 audit ledger; React triage console + approval/audit APIs.
- Result: 205 offline unit/integration tests; 100% precision / 100% recall on held-out **synthetic** abuse, 99.4% recall over 140 randomized runs, zero false-positive blocks on 972 household edge cases.

Aether — Rust Heterogeneous Compute Runtime & LLM Inference Engine

- Built DAG compute runtime with reverse-mode autograd, optimization / kernel-fusion passes, and CPU/GPU execution via WGPU/Metal.
- Implemented GGUF loading, quantized matmul, GQA attention, KV-caching, and streaming inference behind an OpenAI-compatible API with rate limiting and Prometheus metrics; 70+ tests with CPU/GPU cosine-similarity checks.

IronFlow — Rust Workflow Orchestrator

- Built DAG scheduler with concurrent execution, SQLite/WAL persistence, retries, REST APIs, and embedded dashboard; 50-task sequential pipeline at ~85 ms / 12 MB peak per published benchmark methodology.

Additional: Ferrous Browser — async Rust Chrome DevTools Protocol client (Tokio) as lightweight alternative to Node tooling.

TECHNICAL SKILLS

Languages: Rust, Python, JavaScript/TypeScript • **AI/ML:** LLM inference & serving, autograd, agent workflows • **Backend:** Axum, REST, MCP, SQLite, PostgreSQL, NATS • **Systems:** async/concurrency, DAGs, orchestration, GPU compute • **Tools:** Git, Docker, Prometheus